

Predicting Customer Lifetime Value(CLV)

```
In [189]: import pandas as pd #importing Libraries
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [190]: cust_df = pd.read_csv('Customer-Value-Analysis.csv') #Loading dataset into pandas
cust_df.shape
```


Out[190]: (9134, 24)

```
In [191]: cust_df.head()
```

Out[191]:

	Customer	State	Customer Lifetime Value	Response	Coverage	Education	Effective To Date	EmploymentSta
0	BU79786	Washington	2763.519279	No	Basic	Bachelor	2/24/11	Emplo
1	QZ44356	Arizona	6979.535903	No	Extended	Bachelor	1/31/11	Unemplo
2	AI49188	Nevada	12887.431650	No	Premium	Bachelor	2/19/11	Emplo
3	WW63253	California	7645.861827	No	Basic	Bachelor	1/20/11	Unemplo
4	HB64268	Washington	2813.692575	No	Basic	Bachelor	2/3/11	Emplo

5 rows × 24 columns



```
In [192]: cust_df.dtypes #understanding the datatypes
```

```
Out[192]: Customer          object
State          object
Customer Lifetime Value    float64
Response        object
Coverage        object
Education        object
Effective To Date    object
EmploymentStatus    object
Gender          object
Income          int64
Location Code    object
Marital Status    object
Monthly Premium Auto    int64
Months Since Last Claim    int64
Months Since Policy Inception    int64
Number of Open Complaints    int64
Number of Policies    int64
Policy Type        object
Policy          object
Renew Offer Type    object
Sales Channel      object
Total Claim Amount    float64
Vehicle Class      object
Vehicle Size        object
dtype: object
```

```
In [193]: cust_df.isnull().sum() #nullcheck
```

```
Out[193]: Customer          0
State          0
Customer Lifetime Value    0
Response        0
Coverage        0
Education        0
Effective To Date    0
EmploymentStatus    0
Gender          0
Income          0
Location Code    0
Marital Status    0
Monthly Premium Auto    0
Months Since Last Claim    0
Months Since Policy Inception    0
Number of Open Complaints    0
Number of Policies    0
Policy Type        0
Policy          0
Renew Offer Type    0
Sales Channel      0
Total Claim Amount    0
Vehicle Class      0
Vehicle Size        0
dtype: int64
```

Changing Response attribute to 0's and 1's

```
In [194]: cust_df.Response = cust_df.Response.apply(lambda X : 0 if X == 'No' else 1)
          #cust_df.head()
```

Finding correlation between columns

```
In [195]: cust_df.corr()
```

Out[195]:

	Customer Lifetime Value	Response	Income	Monthly Premium Auto	Months Since Last Claim	Months Since Policy Inception	Number of Open Complaints	Number of Policies
Customer Lifetime Value	1.000000	-0.008930	0.024366	0.396262	0.011517	0.009418	-0.036343	0.021955
Response	-0.008930	1.000000	0.011932	0.010966	-0.016597	0.002952	-0.009881	-0.020891
Income	0.024366	0.011932	1.000000	-0.016665	-0.026715	-0.000875	0.006408	-0.008656
Monthly Premium Auto	0.396262	0.010966	-0.016665	1.000000	0.005026	0.020257	-0.013122	-0.011233
Months Since Last Claim	0.011517	-0.016597	-0.026715	0.005026	1.000000	-0.042959	0.005354	0.009136
Months Since Policy Inception	0.009418	0.002952	-0.000875	0.020257	-0.042959	1.000000	-0.001158	-0.013333
Number of Open Complaints	-0.036343	-0.009881	0.006408	-0.013122	0.005354	-0.001158	1.000000	0.001498
Number of Policies	0.021955	-0.020891	-0.008656	-0.011233	0.009136	-0.013333	0.001498	1.000000
Total Claim Amount	0.226451	0.016877	-0.355254	0.632017	0.007563	0.003335	-0.014241	-0.002354


```
In [196]: corr_matrix = cust_df.corr()['Customer Lifetime Value'] #sorting correlation  
sorted_corr = corr_matrix.sort_values(ascending = False)  
sorted_corr
```

```
Out[196]: Customer Lifetime Value      1.000000  
Monthly Premium Auto      0.396262  
Total Claim Amount      0.226451  
Income      0.024366  
Number of Policies      0.021955  
Months Since Last Claim      0.011517  
Months Since Policy Inception      0.009418  
Response      -0.008930  
Number of Open Complaints      -0.036343  
Name: Customer Lifetime Value, dtype: float64
```

Also we can visualize this using heat map

```
In [197]: sns.set(rc={'figure.figsize':(10,10)}) #heatmap
sns.heatmap(cust_df.corr(),annot = True)
```

```
Out[197]: <matplotlib.axes._subplots.AxesSubplot at 0x137e2fa20f0>
```




As we can see that the monthly premium auto, total claim amount, and income has an effect on customer lifetime value attribute. Let's work on these attributes

Outliers check


```
In [198]: sns.set(rc={'figure.figsize':(8,6)})  
x = cust_df['Income'].values  
ax = sns.boxplot(x)  
print('The meadian is: ', cust_df['Income'].median())
```

The meadian is: 33889.5




```
In [199]: sns.set(rc={'figure.figsize':(8,6)})  
x = cust_df['Monthly Premium Auto'].values  
ax = sns.boxplot(x)  
print('The meadian is: ', cust_df['Monthly Premium Auto'].median())
```

The meadian is: 83.0



```
In [200]: sns.set(rc={'figure.figsize':(8,6)})  
x = cust_df['Total Claim Amount'].values  
ax = sns.boxplot(x)  
print('The meadian is: ', cust_df['Total Claim Amount'].median())
```

The meadian is: 383.94543350000004




With results above, we can say that there are outliers in total claim amount and monthly premium auto. Usually we remove the outliers but as this data is related to insurance company, the outliers can be our most important customers. So, removing the outliers is not needed.


```
In [201]: #normal distributions
sns.set(rc = {'figure.figsize': (6,6)})
sns.distplot(cust_df['Income'])
mean = cust_df['Income'].mean()
plt.axvline(mean,0,1,color = 'red')
```

C:\ProgramData\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

Out[201]: <matplotlib.lines.Line2D at 0x137e2094ef0>



```

In [202]: fig, axes = plt.subplots(2, 2, figsize=(12, 7))

a = cust_df['Income'].values
b = cust_df['Monthly Premium Auto'].values
c = cust_df['Total Claim Amount'].values
d = cust_df['Number of Policies'].values

# plot 1
sns.distplot(a, color = 'blue', ax=axes[0,0])


# plot 2
sns.distplot(b, color = 'blue', ax=axes[0,1])

# plot 3
sns.distplot(c, color = 'blue', ax=axes[1,0])

#plot 4
sns.distplot(d, color = 'blue', ax=axes[1,1])

```

Out[202]: <matplotlib.axes._subplots.AxesSubplot at 0x137df4b83c8>



As the columns are skewed, let's try normalizing it by applying transformations such as square

```

In [203]: fig, axes = plt.subplots(2, 2, figsize=(12, 7))
          at = (cust_df['Income']**2).values
          bt = (cust_df['Monthly Premium Auto']**2).values
          ct = (cust_df['Total Claim Amount']**2).values
          dt = (cust_df['Number of Policies']**2).values

          # plot 1
          sns.distplot(at, color = 'blue', ax=axes[0,0])


          # plot 2
          sns.distplot(bt, color = 'blue', ax=axes[0,1])

          # plot 3
          sns.distplot(ct, color = 'blue', ax=axes[1,0])

          # plot 4
          sns.distplot(dt, color = 'blue', ax=axes[1,1])

```

Out[203]: <matplotlib.axes._subplots.AxesSubplot at 0x137e3b9dc88>



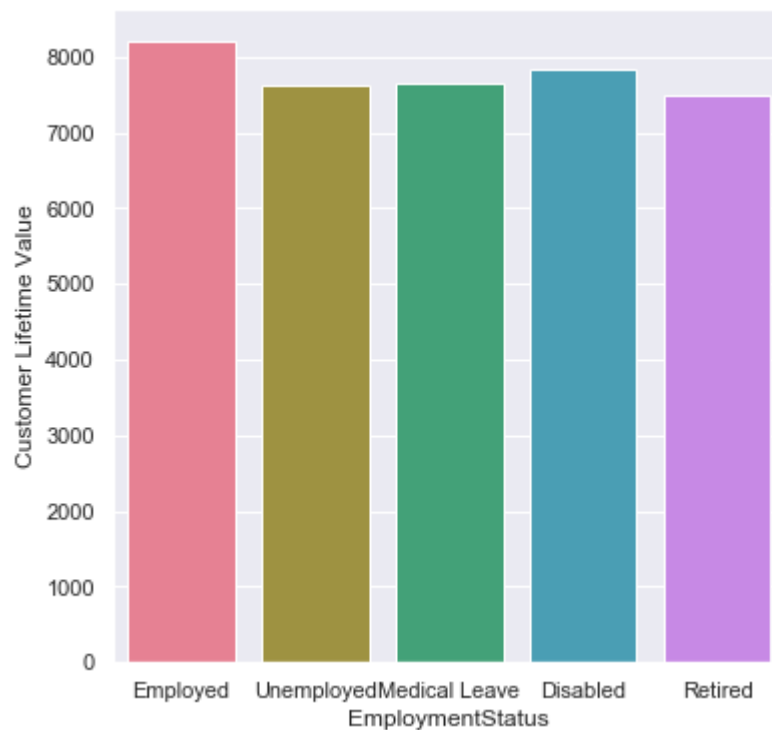
With the transformation of the columns, our data is getting much skewed or the peaks of the distribution are increasing. So, it is better to remain with our original data

visualizations

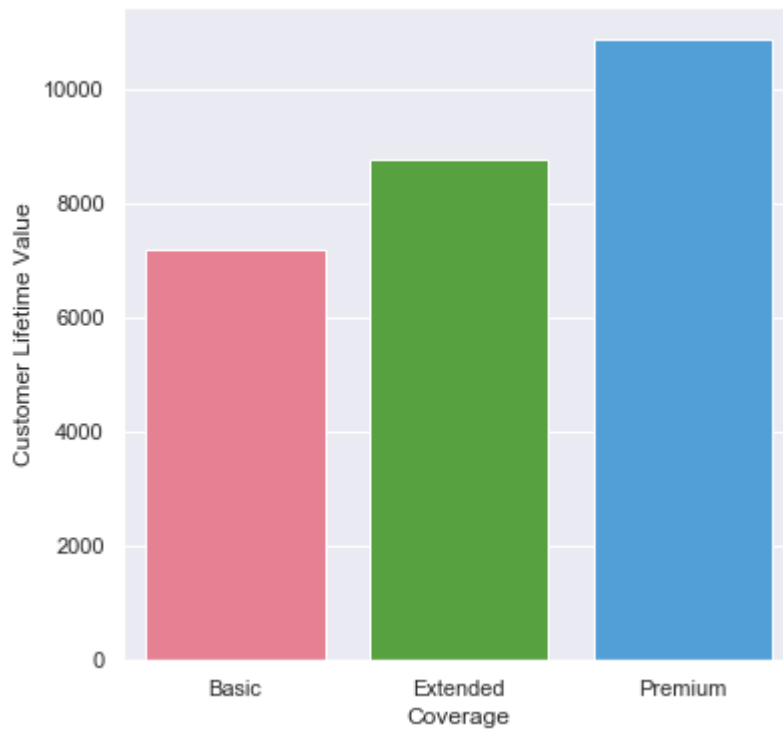
As it is important for the business to understand their customers behavior. some key points can be drawn if we observe the trends in response attribute w.r.t other columns. (Tableau)

Analysis on customer life time value(categorical variables)

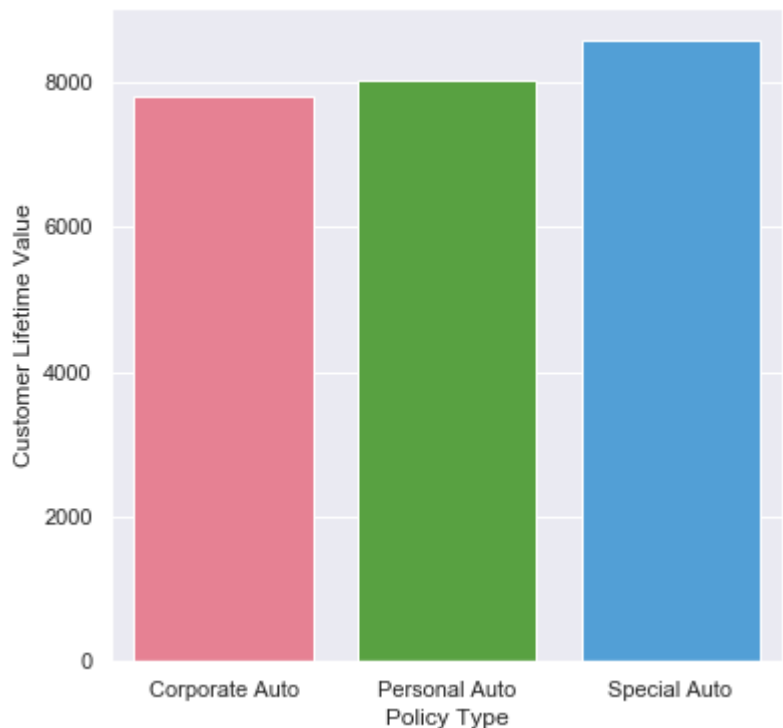
```
In [204]: # How is customer lifetime value related to employment status of thr customer  
ax = sns.barplot(y='Customer Lifetime Value', x='EmploymentStatus', data = cust_d
```



```
In [205]: #CLV w.r.t coverage  
ax = sns.barplot(y='Customer Lifetime Value', x='Coverage', data = cust_df, ci = 1)
```




```
In [206]: #CLV w.r.t policy type  
ax = sns.barplot(y='Customer Lifetime Value', x='Policy Type', data = cust_df, ci = 1)
```



As there is no much variaion in CLV w.r.t whether a customer is employed or not, we can disregard this column.

```
In [207]: #CLV w.r.t gender  
bx = sns.barplot(y='Customer Lifetime Value', x='Gender',data = cust_df, ci = Fa
```



```
In [208]: #CLV w.r.t Education and Gender  
ax = sns.lineplot(y='Customer Lifetime Value', x='Education',hue = 'Gender',data
```



We can say that people who are educated(doctor) have much less customer lifetime value, when compared to the people who studied high school or below. And also, CLV is slightly high in case of

females, than males.

```
In [209]: rough = cust_df.drop(['State', 'Customer', 'Response', 'EmploymentStatus', 'Gender'],
#cust_df.columns
```

```
In [210]: rough_cat = cust_df.select_dtypes(include = ['object']).columns
rough_cat
```

```
Out[210]: Index(['Coverage', 'Marital Status', 'Renew Offer Type', 'Vehicle Class'], dtype
e='object')
```

```
In [211]: cols = ['Coverage', 'Marital Status', 'Renew Offer Type', 'Vehicle Class'] #dummy
new = pd.get_dummies(cust_df, columns=['Coverage', 'Marital Status', 'Number of Pol:
```

```
In [212]: new.dtypes
```

```
Out[212]: Customer Lifetime Value      float64
Monthly Premium Auto                  int64
Months Since Last Claim               int64
Months Since Policy Inception         int64
Number of Open Complaints             int64
Total Claim Amount                   float64
Coverage_Extended                    uint8
Coverage_Premium                     uint8
Marital Status_Married                uint8
Marital Status_Single                 uint8
Number of Policies_2                  uint8
Number of Policies_3                  uint8
Number of Policies_4                  uint8
Number of Policies_5                  uint8
Number of Policies_6                  uint8
Number of Policies_7                  uint8
Number of Policies_8                  uint8
Number of Policies_9                  uint8
Renew Offer Type_Offer2               uint8
Renew Offer Type_Offer3               uint8
Renew Offer Type_Offer4               uint8
Vehicle Class_Luxury Car              uint8
Vehicle Class_Luxury SUV              uint8
Vehicle Class_SUV                     uint8
Vehicle Class_Sports Car              uint8
Vehicle Class_Two-Door Car            uint8
dtype: object
```

In [213]: `new.head()`

Out[213]:

	Customer Lifetime Value	Monthly Premium Auto	Months Since Last Claim	Months Since Policy Inception	Number of Open Complaints	Total Claim Amount	Coverage_Extended	Cover
0	2763.519279	69	32	5	0	384.811147	0	
1	6979.535903	94	13	42	0	1131.464935	1	
2	12887.431650	108	18	38	0	566.472247	0	
3	7645.861827	106	18	65	0	529.881344	0	
4	2813.692575	73	12	44	0	138.130879	0	

5 rows × 26 columns



Model Selection

In [214]: `X = new.drop(['Customer Lifetime Value'],axis = 1)`
`#X.head()`
`y = new[['Customer Lifetime Value']]`
`#print(y)`

In [215]: `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)`

In [219]: `from sklearn.linear_model import LinearRegression`
`import sklearn.metrics as sm`
`regressor = LinearRegression()`
`regressor.fit(X_train, y_train)`
`y_pred = regressor.predict(X_test)`
`print(regressor.score(X_test, y_test))`
`print("R2 score =", round(sm.r2_score(y_test, y_pred), 2))`

0.6335667047027755

R2 score = 0.63


```
In [217]: from sklearn.tree import DecisionTreeRegressor
regr_1 = DecisionTreeRegressor(max_depth = 5, random_state=1)
regr_1.fit(X_train, y_train)
ypred_dt = regr_1.predict(X_test)
print(regr_1.score(X_test, y_test))
print("R2 score =", round(sm.r2_score(y_test, ypred_dt), 2))
```

0.6177902611729776

R2 score = 0.62